

*#This is a Student Performance Dataset designed to examine the factors influencing academic student performance.
#In this notebook, I have implemented multiple linear regression from scratch.*

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
import seaborn as sns
from sklearn import linear_model
import matplotlib.pyplot as plt
```

```
data=pd.read_csv(r"C:\Users\shamzkha\Documents\
Student_Performance.csv")
```

```
data.shape
```

```
(10000, 6)
```

```
data.head()
```

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep
0	7	99	Yes	
1	4	82	No	
2	8	51	Yes	
3	5	52	Yes	
4	7	75	No	

	Sample Question Papers Practiced	Performance Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0

```
data.isnull().sum()
```

Hours Studied	0
Previous Scores	0
Extracurricular Activities	0
Sleep Hours	0
Sample Question Papers Practiced	0
Performance Index	0

dtype: int64

```

#No missing values found

#Renaming the columns

data=data.rename(columns={"Hours Studied":"Hours_Studied"})
data=data.rename(columns={"Previous Scores":"Previous_Scores"})
data=data.rename(columns={"Extracurricular
Activities":"Extracurricular_Activities"})
data=data.rename(columns={"Sleep Hours":"Sleep_Hours"})
data=data.rename(columns={"Sample Question Papers Practiced":"SQPP"})
data=data.rename(columns={"Performance Index":"Performance_Index"})

data.dtypes

Hours_Studied                int64
Previous_Scores              int64
Extracurricular_Activities   object
Sleep_Hours                  int64
SQPP                         int64
Performance_Index            float64
dtype: object

```

Converting categorical value to numerical value

```

data.Extracurricular_Activities=le.fit_transform(data.Extracurricular_
Activities)

#data.dtypes

```

EDA

```

#The target column performance index is numerical so no class
imbalance treatment is required

```


Outlier Treatment

```

sns.boxplot(data=data,y="Hours_Studied")


<Axes: ylabel='Hours_Studied'>

```




```
sns.boxplot(data=data,y="Previous_Scores")
```


```
<Axes: ylabel='Previous_Scores'>
```



```
sns.boxplot(data=data,y="Sleep_Hours")  
<Axes: ylabel='Sleep_Hours'>
```



```
sns.boxplot(data=data,y="SQPP")  
<Axes: ylabel='SQPP'>
```



#extracurricular activities not checked as it is categorical column
#No outlier ,so no treatment is required

Correlation

```
data1=data.corr()
data1
```

	Hours_Studied	Previous_Scores	\
Hours_Studied	1.000000	-0.012390	
Previous_Scores	-0.012390	1.000000	
Extracurricular_Activities	0.003873	0.008369	
Sleep_Hours	0.001245	0.005944	
SQPP	0.017463	0.007888	
Performance_Index	0.373730	0.915189	
	Extracurricular_Activities	Sleep_Hours	
SQPP \			
Hours_Studied	0.003873	0.001245	
0.017463			
Previous_Scores	0.008369	0.005944	
0.007888			
Extracurricular_Activities	1.000000	-0.023284	
0.013103			

Sleep_Hours	-0.023284	1.000000
0.003990		
SQPP	0.013103	0.003990
1.000000		
Performance_Index	0.024525	0.048106
0.043268		

	Performance_Index
Hours_Studied	0.373730
Previous_Scores	0.915189
Extracurricular_Activities	0.024525
Sleep_Hours	0.048106
SQPP	0.043268
Performance_Index	1.000000

#Conclusion:

#Performance Index is most strongly correlated with Previous Scores (0.915).

#Hours Studied has a moderate, positive correlation with Performance Index (0.374).

#Extracurricular Activities, Sleep Hours, and SQPP show weak correlations with performance.

Linear Regression

```
data.head()
```

	Hours_Studied	Previous_Scores	Extracurricular_Activities
Sleep_Hours \			
0	7	99	1
9			
1	4	82	0
4			
2	8	51	1
7			
3	5	52	1
5			
4	7	75	0
8			

	SQPP	Performance_Index
0	1	91.0
1	2	65.0
2	2	45.0
3	2	36.0
4	5	66.0

```

#Separating features from target
x=data.iloc[:,0:5]
#x.head()

y=data.iloc[:,5]
#y.head()

import sklearn
from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=101)

x_train.shape,x_test.shape,y_train.shape,y_test.shape
((8000, 5), (2000, 5), (8000,), (2000,))

from sklearn import linear_model
linear=linear_model.LinearRegression()
linear.fit(x_train,y_train)
LinearRegression()
pred=linear.predict(x_test)
pred
array([44.38891422, 96.15853705, 30.52422606, ..., 45.9649391 ,
       59.64558417, 16.80314704])

linear.coef_
array([2.85283863, 1.01817717, 0.63290609, 0.48524068, 0.19369907])

linear.intercept_
-34.10276056362111

R2=linear.score(x_train,y_train) #Rsquare value
R2
0.9887109739552409

Adj_R2=1-(((1-R2)*(8000-1))/(8000-5-1)) #adjusted Rsquare value
Adj_R2
0.9887039130182602

pred_train=linear.predict(x_train)
#pred_train
pred_train.shape

```


```

(8000,)
mean_y=y_train.mean()
mean_y
55.260625
SSE=np.sum(np.square(pred_train-y_train))
SSE
33333.5316344614
SSR=np.sum(np.square(pred_train-mean_y))
SSR
2919404.0652405405
Rsqr=SSR/(SSR+SSE) #Rsquare value by formula
Rsqr
0.9887109739552408
from sklearn import metrics
#MAE=mean absolute error
MSE=metrics.mean_squared_error(pred,y_test)
MSE
4.091042932500655
RMSE=np.sqrt(MSE)
RMSE
2.0226326736460716
#MAPE=Mean absolute Percentage Error
error=pred-y_test
error
error_abs=np.abs(error)
#error_abs
MAPE=np.mean(error_abs/y_test)*100
MAPE
3.512131434073932
Accuracy=(100-MAPE)
Accuracy
96.48786856592606

```




```
plt.scatter(pred,y_test)
plt.show()
```



```
from scipy import stats
slope,intercepts,r,p,std_err=stats.linregress(pred,y_test)
def myfunc(y_test):
    return slope*y_test+intercepts
mymodel=list(map(myfunc,y_test))

plt.scatter(pred,y_test)
plt.plot(y_test,mymodel)
plt.show()
```



#Conclusion:

#The linear regression model developed for predicting the dependent variable has demonstrated strong predictive performance, with an accuracy of 96.48%.

#This indicates that the model is highly effective at making predictions, with minimal error.

#The adjusted R-squared value of 0.988870 suggests that approximately 98.89% of the variance in the dependent variable can be explained by the independent variables in the model.

L1=Lasso

```
from sklearn.linear_model import Lasso
lasso=Lasso()

lasso.fit(x_train,y_train)

Lasso()

lasso.coef_

array([2.71029613, 1.01484602, 0.          , 0.13493396, 0.07585949])
```

```

l1_pred=lasso.predict(x_test)
l1_pred

array([45.53780956, 95.07744207, 31.42387377, ..., 45.41309159,
       60.24373499, 16.69862985])

l1_R2=lasso.score(x_train,y_train)
l1_R2

0.9867853232610752

l1_adj_R2=1-(((1-l1_R2)*(8000-1))/(8000-5-1))
l1_adj_R2

0.9867770578890843

df=pd.DataFrame({"Feature_importances":lasso.coef_,"columns":list(x)})
df

```

	Feature_importances	columns
0	2.710296	Hours_Studied
1	1.014846	Previous_Scores
2	0.000000	Extracurricular_Activities
3	0.134934	Sleep_Hours
4	0.075859	SQPP

```

df2=pd.DataFrame({"Actual":y_test,"Predictions":l1_pred})
df2

```

	Actual	Predictions
6676	43.0	45.537810
6421	95.0	95.077442
9834	29.0	31.423874
8492	48.0	51.069095
9982	44.0	44.171937
...
4441	38.0	39.997799
4166	42.0	39.652071
2567	46.0	45.413092
8527	61.0	60.243735
406	17.0	16.698630

```

[2000 rows x 2 columns]


MSE_l1=metrics.mean_squared_error(l1_pred,y_test)
MSE_l1

4.718953466209578

sns.lmplot(x="Actual",y="Predictions",data=df2,fit_reg=False)
d_line=np.arange(df2.min().min(),df2.max().max())

```

```
plt.plot(d_line,color="red",linestyle="-")
plt.show()
```



L2=Ridge

```
from sklearn.linear_model import Ridge
rd=Ridge()

rd.fit(x_train,y_train)

Ridge()

rd_pred=rd.predict(x_test)
rd_pred

array([44.38903204, 96.15814176, 30.52429609, ..., 45.96509303,
       59.64559074, 16.80317541])

list(rd.coef_)
```

```
[2.8527862054001183,  
1.0181766862753776,  
0.6325885585926317,  
0.4852177273122147,  
0.1936980544421156]
```

```
rd_R2=rd.score(x_train,y_train)  
rd_R2
```

```
0.9887109738325554
```

```
rd_adj_R2=1-(((1-rd_R2)*(8000-1))/(8000-5-1))  
rd_adj_R2
```

```
0.9887039128954979
```

```
df_1=pd.DataFrame({"Feature_importances":rd.coef_,"columns":list(x)})  
df_1
```

	Feature_importances	columns
0	2.852786	Hours_Studied
1	1.018177	Previous_Scores
2	0.632589	Extracurricular_Activities
3	0.485218	Sleep_Hours
4	0.193698	SQPP

```
df_2=pd.DataFrame({"Actual":y_test,"Predictions":rd_pred})  
df_2
```

	Actual	Predictions
6676	43.0	44.389032
6421	95.0	96.158142
9834	29.0	30.524296
8492	48.0	49.753932
9982	44.0	43.266933
...
4441	38.0	39.831898
4166	42.0	38.667764
2567	46.0	45.965093
8527	61.0	59.645591
406	17.0	16.803175


```
[2000 rows x 2 columns]
```

```
MSE_rd=metrics.mean_squared_error(rd_pred,y_test)  
MSE_rd
```

```
4.0910252511709615
```

```
sns.lmplot(x="Actual",y="Predictions",data=df_2,fit_reg=False)  
d_line=np.arange(df_2.min().min(),df_2.max().max())
```

```
plt.plot(d_line,color="red",linestyle="-")
plt.show()
```



Feature Selection

```
df
```

	Feature_importances	columns
0	2.710296	Hours_Studied
1	1.014846	Previous_Scores
2	0.000000	Extracurricular_Activities
3	0.134934	Sleep_Hours
4	0.075859	SQPP

```
data.shape
```

```
(10000, 6)
```

```
l_new=data.drop(["Extracurricular_Activities","SQPP"],axis=1)
```

```
#dropping the non-significant columns
```

```
l_new.shape
```

```
(10000, 4)
```

```
l_new.head()
```

	Hours_Studied	Previous_Scores	Sleep_Hours	Performance_Index
0	7	99	9	91.0
1	4	82	4	65.0
2	8	51	7	45.0
3	5	52	5	36.0
4	7	75	8	66.0

```
l_new.isnull().sum()
```

```
Hours_Studied      0
Previous_Scores    0
Sleep_Hours        0
Performance_Index  0
dtype: int64
```

```
l_new.dtypes
```

```
Hours_Studied      int64
Previous_Scores    int64
Sleep_Hours        int64
Performance_Index  float64
dtype: object
```

```
x1=l_new.iloc[:,0:3] #separating significant features from target
```

```
x1.head()
```

	Hours_Studied	Previous_Scores	Sleep_Hours
0	7	99	9
1	4	82	4
2	8	51	7
3	5	52	5
4	7	75	8

```
y1=l_new.iloc[:,3]
```

```
y1.head()
```

```
0    91.0
1    65.0
2    45.0
3    36.0
4    66.0
```

```
Name: Performance_Index, dtype: float64
```

```

x1_train,x1_test,y1_train,y1_test=train_test_split(x1,y1,test_size=0.2
,random_state=101)
x1_train.shape,x1_test.shape,y1_train.shape,y1_test.shape

((8000, 3), (2000, 3), (8000,), (2000,))

linear.fit(x1_train,y1_train)

LinearRegression()

linear_new_pred=linear.predict(x1_test)
linear_new_pred

array([45.03684486, 95.97802431, 30.87754813, ..., 45.97886901,
       59.91952586, 16.5575236 ])

new_R2=linear.score(x1_train,y1_train)
new_R2

0.9875902655142141

Anew_R2=1-(((1-new_R2)*(8000-1))/(8000-3-1))
Anew_R2

0.9875856095357928

new_MSE=metrics.mean_squared_error(linear_new_pred,y1_test)
new_MSE

4.470622176935699

df_new=pd.DataFrame({"Actual_n":y1_test,"Predictions_n":linear_new_pre
d})
df_new

   Actual_n  Predictions_n
6676     43.0      45.036845
6421     95.0      95.978024
9834     29.0      30.877548
8492     48.0      50.374289
9982     44.0      44.299792
...      ...      ...
4441     38.0      39.024597
4166     42.0      38.068905
2567     46.0      45.978869
8527     61.0      59.919526
406      17.0      16.557524


[2000 rows x 2 columns]

sns.lmplot(x="Actual_n",y="Predictions_n",data=df_new,fit_reg=False)
d_line=np.arange(df_new.min().min(),df_new.max().max())

```



```
plt.plot(d_line,color="red",linestyle="-")
plt.show()
```



```
list1=["Linear Regression","Lasso","Ridge","Feature_SelectionModel"]
list2=[R2,l1_R2,rd_R2,new_R2]
list3=[Adj_R2,l1_adj_R2,rd_adj_R2,Anew_R2]
list4=[MSE,MSE_l1,MSE_rd,new_MSE]

Final_Result=pd.DataFrame({"Model_Name":list1,"R2_value":list2,"Adj_R2":list3,"MSE":list4})
Final_Result
```

	Model_Name	R2_value	Adj_R2	MSE
0	Linear Regression	0.988711	0.988704	4.091043
1	Lasso	0.986785	0.986777	4.718953
2	Ridge	0.988711	0.988704	4.091025
3	Feature_SelectionModel	0.987590	0.987586	4.470622

#Conclusion:

#The Linear Regression and Ridge models are the best, showing

identical Rsquare, adjusted Rsquare, and MSE values.
#In comparison, the Lasso and Feature Selection Model show slightly lower performance with R-squared values of 0.986785 and 0.987590, as well as higher MSEs.
#In summary, the high accuracy and adjusted R-squared value confirm that the linear regression model and Ridge is both reliable and capable of providing meaningful insights for prediction.